

What AI-native could mean for DFRNT and Urgent

A strategic discussion deck to help the team understand the future-state vision, the path to get there, and the three real options in front of us.

Prepared from the current AI-native strategy pack, white paper, and latest strategic framing update.

This is not a “how do we use more AI?” conversation

THE WRONG FRAME

- Add copilots to current workflows
- Automate a few admin tasks
- Project today’s org chart and systems into the future
- Optimise the current courier model at the edges

THE RIGHT FRAME

- Start with a compelling 5–7 year vision
- Ask what an AI-native operating model actually looks like
- Backcast the systems, workflows and capabilities needed
- Decide whether we are building a better old model — or a new one

We are uniquely positioned to do this

URGENT

Real operator

Live dispatch complexity, real customers, real margin pressure, real service failures, real proof opportunities.

DFRNT

Platform builder

The software company that can encode the learning, automate the workflows, and turn operational insight into scalable product capability.

No other courier company is better positioned than we are to define what an AI-native transport business could actually look like in practice.

What are we backcasting from?

If we designed DFRNT and Urgent from scratch today to be AI-native 5–7 years from now, what would exist, what would be automated, what would humans still own, and what would make the business structurally better than a traditional courier operator?

Continuously learning operating system

Operational events are captured once and reused everywhere. The platform gets smarter from real execution, not reports written after the fact.

Exception-led human work

Agents, rules and workflows handle repetitive coordination. Humans focus on exceptions, judgement, service design and commercial decisions.

Better scaling economics

Quality, speed and margin improve without management overhead scaling at the same rate.

AI-native is not “chatbot everywhere”

Visible

The data exists and can be seen.



Assisted

AI drafts, suggests, classifies and summarises.



Automated

Low-risk work completes end-to-end through rules and agents.

Core design idea

- Every stage has a clean system of record
- Key events are emitted once and reused everywhere
- Apps are organised around operating outcomes, not historical team boundaries

End-state discipline

- Humans stay on judgement, exceptions and relationships
- Machines take coordination, classification and repetitive decision support
- The system reprioritises around cost, SLA risk and customer value

The four operating stages

STAGE 01

Sales

Win the customer, price the work, configure the service, and hand off cleanly into live operations.

STAGE 02

Operations

Turn demand into planned, controlled, profitable execution through orchestration and exception control.

STAGE 03

Accounts

Bill accurately, settle correctly, reconcile fast, and surface margin truth.

STAGE 04

Delivery

Execute the work in the field, capture proof, exceptions and customer visibility.

Lead → Quote → Onboard → Book → Plan → Dispatch → Deliver → Invoice
→ Settle → Report → Improve

Why this matters economically

CUSTOMER ECONOMICS

Lower cost to serve

AI-native workflow redesign can change the cost structure of delivery rather than just improving software around the same manual process.

OPERATING LEVERAGE

20 staff → 5

If the operating model is redesigned around automation and exception handling, staffing intensity can fall materially over time.

STRATEGIC UPSIDE

Better proof

Urgent becomes the proving ground and DFRNT becomes the product engine — creating a stronger valuation and market-entry story.

There are three credible paths

Option 1

Stay AI-enabled

- Add selective AI features to the current model
- Lowest near-term disruption
- Highest risk of becoming a cleaner version of the old model

Option 2

Dual-track: defend + wedge

- Usually sounds balanced
- Works only if “defend the base” does not consume all oxygen
- High risk of delaying the real pivot

Option 3

Full AI-native pivot

- Make the pivot the primary build priority
- Use Urgent as the proof case
- Best fit if the aim is structural change, not incremental optimisation

What each option gives up

Option 1 risk

We end up with better tooling, but not a structurally better business.

Option 2 risk

Legacy work absorbs the team and the wedge never becomes the strategy.

Option 3 risk

Requires discipline, proof milestones, and the willingness to stop saying yes to work that drags us backwards.

BIGGEST EXTERNAL RISK

Prospects keep asking DFRNT to preserve manual process steps that an AI-native strategy is supposed to eliminate.

BIGGEST INTERNAL RISK

We confuse AI novelty with operating-model change and call that progress too early.

What has to be true if we want to win

- Urgent is treated as the proof case, not just another tenant
- The roadmap is ranked against the target operating model
- Build capacity for the AI-native wedge is protected
- New work is judged on whether it moves us toward the future-state vision

- We measure proof in real operating metrics
- We keep the system-of-record and event model clean
- We do not let tenant-by-tenant requests define the product architecture
- We build a testimony engine from the transformation story

RECOMMENDATION

Start with the vision. Backcast the path. Bias toward the pivot.

The strongest direction is a more decisive AI-native pivot — but grounded in proof milestones, real operating metrics, and roadmap discipline.



Questions for the team

- What 5–7 year future state are we actually aiming for?
- Which of the three options are we really choosing?
- What are we explicitly not choosing?

- What proof do we need in the next 12 months?
- What tenant or roadmap work should we stop doing if it pulls us backward?
- What has to be protected so this does not become just another side project?